

SHIELDMCP: An Interactive System for Runtime Security of Model Context Protocol Tool Calls

Saurabh Yergattikar

Independent Researcher

Dublin, CA, USA

saurabh.ssy@gmail.com

Abstract

The Model Context Protocol (MCP) has become the dominant standard for connecting large language models to external tools, but it exposes a new attack surface: an adversary can compromise an agent not through the user prompt, but through the *tools* the agent trusts—poisoning tool descriptions, injecting hidden instructions into tool responses, or distributing malicious servers. We present SHIELDMCP, an interactive, deployment-ready system that secures the MCP tool layer at runtime. SHIELDMCP operates as a transparent proxy between the LLM client and any MCP server, performing three-stage validation—*pre-call* description integrity, *outbound* parameter sanitization, and *post-response* analysis—without modifying either the agent or the server. The demonstration exposes the full pipeline through a browser interface in which visitors launch curated attacks from five threat families, watch each validation stage execute live with per-token highlighting of the detected payload, and submit their own tool descriptions, call parameters, or tool responses to a playground that runs the real detector. On a benchmark of 487 attack and 200 benign scenarios across 40 MCP servers, SHIELDMCP reduces the attack success rate from 70.9% (undefended) to 8.6% with rule-based detectors and to near-zero with optional fine-tuned classifiers, while preserving 95% of benign task utility and adding under 20 ms of median latency per tool call. SHIELDMCP is open-source under the Apache-2.0 license and installable with a single `pip` command.

1 Introduction

Agents built on large language models (LLMs) increasingly act on the world by calling external tools: reading files, querying databases, searching the web, and executing code. The Model Context Protocol (MCP; [Salih et al., 2025](#)), now governed by the Linux Foundation and adopted by An-

thropic, OpenAI, and Google, has become the predominant standard for this tool-calling capability, with over 16,000 publicly available MCP servers as of early 2026.

This convenience introduces an attack surface that existing LLM safety measures do not cover. Traditional guardrails inspect the *user-facing* channel—prompts in and completions out. MCP-enabled agents face a qualitatively different risk: the adversary need not interact with the user at all, and can instead compromise the *tool layer* the agent implicitly trusts. Poisoned tool descriptions steer an agent’s planning ([Wang et al., 2025](#)); hidden instructions embedded in tool *responses* hijack subsequent actions ([Greshake et al., 2023](#); [An et al., 2025](#)); and malicious servers reach public registries with minimal vetting ([Song et al., 2025](#); [Zhao et al., 2025](#)). These are not hypothetical: CVE-2025-6514 compromised over 437,000 developer environments through a malicious MCP package, and empirical studies report tool-poisoning success rates of 61–73% against production agents ([Wang et al., 2025](#)).

A companion research paper introduces the threat taxonomy and defense framework that SHIELDMCP implements. This paper describes the *system and its interactive demonstration*: a working, open-source, installable proxy that makes MCP tool-layer attacks—and their runtime mitigation—visible and reproducible. Our contributions as a demonstration are: (i) a transparent, drop-in MCP security proxy that requires no changes to the agent or server and supports both stdio and SSE transports; (ii) a browser-based interface that executes the real detection pipeline live, animating each of the three validation stages and highlighting the exact tokens that trigger a block; and (iii) an interactive playground and reproducible benchmark harness (487 attacks, 200 benign tasks, 40 servers) that lets visitors probe the system with their own inputs.

2 The Problem and Target Audience

What problem does the system address?

SHIELD MCP defends the MCP tool layer against five attack families that production deployments must confront: *tool poisoning* via metadata manipulation, *indirect prompt injection* (IPI) through tool responses, *supply-chain* compromise of MCP servers, *rug pulls* in which a trusted tool silently mutates its behavior, and *cross-tool exploitation chains* that compose benign calls into an exfiltration pipeline. Each of these bypasses input/output guardrails because the malicious content rides on the tool channel, which those guardrails do not inspect.

Why is it important? As MCP adoption accelerates under Linux Foundation stewardship, the tool layer is becoming critical infrastructure for agentic AI, yet the ecosystem lacks standardized vetting: Zhao et al. (2025) find that 27.2% of analyzed servers expose exploitable tools. In regulated sectors, an unmonitored tool layer is a direct compliance risk, since multi-stage inference attacks can exfiltrate confidential data even when standard safety measures are in place. A *deployable, modification-free* runtime defense is therefore of immediate practical value.

Who is the target audience? The primary audience is engineers and security teams integrating MCP into production agents, who can deploy SHIELD MCP as a proxy in front of untrusted servers. The demonstration also targets NLP researchers studying agent security, who can use the open benchmark and playground to reproduce attacks and evaluate new detectors, and educators, who can use the live interface to teach tool-layer threat models concretely.

3 System Architecture

SHIELD MCP operates as a transparent proxy on the JSON-RPC transport between the LLM client and connected MCP servers (Figure 1). Because it speaks the MCP wire protocol, it is a drop-in interposition: agents send standard requests and servers receive standard requests. A single command wraps any stdio server,

```
shieldmcp proxy - python
server.py
```

and an SSE proxy mode covers HTTP-based servers. The pipeline performs three stages.




Figure 1: SHIELD MCP interposes on the JSON-RPC stream as a transparent proxy, validating tool descriptions before planning, outbound parameters before they leave the client, and inbound responses before the agent reads them.

Stage 1: Description integrity (pre-call). Before the agent’s planning phase, SHIELD MCP inspects every registered tool description with two checks. A *structural* analyzer parses descriptions for hidden Unicode (zero-width joiners, right-to-left overrides), HTML/markdown injection, and instruction-like patterns embedded in metadata fields, catching the “invisible prompt” class. A *semantic* check then scores whether the description contains embedded action directives. SHIELD MCP also maintains a baseline registry of vetted tool signatures (name, description hash, parameter schema) and raises a high-severity alert when a previously seen tool changes—the rug-pull detector—requiring explicit re-authorization.

Stage 2: Parameter sanitization (outbound). When the agent constructs a call, SHIELD MCP validates the outbound parameters against the tool’s declared schema (type, range) and scans string arguments for payloads that a downstream system could interpret: SQL injection fragments, shell command injection, path traversal, and prompt-injection strings targeting other LLMs in the chain.

Stage 3: Response analysis (post-response). Tool responses are the primary vector for indirect prompt injection, so this is the most consequential stage. SHIELD MCP performs (i) *instruction detection*, identifying segments of the response that contain directive language using either pattern rules or a token-level classifier; (ii) *context-boundary enforcement*, wrapping responses in explicit delimiters that the agent is conditioned to treat as untrusted data—a lightweight, prompt-engineering approximation of the control/data separation advocated by Debenedetti et al. (2025); and (iii) *cross-call correlation*, maintaining a

session-level dependency graph so that a response which triggers an unexpected follow-on tool call raises a high-severity chaining alert.

Pluggable detection backends. Stages 1 and 3 each expose a backend switch. The default `heuristic` backend is dependency-light and requires no training or GPU, making it the conservative operating point for deployment. An optional `classifier` backend loads fine-tuned DistilBERT models—a binary sequence classifier for descriptions and a token classifier for responses—for higher recall. Backends are selected in a YAML config or at the command line, so operators trade recall against footprint without code changes.

4 The Interactive Demonstration

The demonstration is a single-page web application, served by the same Python package as the proxy and launched with `shieldmcp demo`. Every verdict shown in the browser is computed live by the real pipeline—nothing is mocked or pre-recorded. The interface (Figure 2) has three regions.

Attack gallery. A left panel lists curated scenarios spanning all five threat families plus a benign control. Selecting one streams the corresponding MCP traffic through the proxy and animates the three stage cards in sequence: each card transitions from *scanning* to a *pass/warn/block* verdict, reports its measured latency and alert count, and—critically for pedagogy—renders the offending tool description, parameter, or response with the exact substrings that triggered detection highlighted in place. When a stage blocks, the pipeline halts there (mirroring production behavior), and a verdict banner summarizes the end-to-end outcome and latency.

Live benchmark. A right panel surfaces the aggregate evaluation—overall attack success rate, benign pass rate, and a per-family attack-success bar chart—read directly from the benchmark result file, so the numbers a visitor sees on screen are the same numbers reported in Section 5.

Playground. A modal lets visitors paste an arbitrary tool description, call parameters, or tool response and run it through all three stages, returning the per-stage verdict, matched attack family, and latency. This turns the demo from a scripted

Attack family	Total	ASR (heur.)	ASR (clf.)
Tool poisoning	120	5.0%	0.0%
IPI via response	110	16.4%	0.0%
Supply chain	90	20.0%	0.0%
Rug pull	75	0.0%	0.0%
Cross-tool chain	92	0.0%	0.0%
Overall	487	8.6%	0.0%

Table 1: Attack success rate (ASR, lower is better) on the framework benchmark with the heuristic and fine-tuned-classifier backends. Undefended ASR averages 70.9% across five LLM backends in the companion study.

tour into an interactive probe: reviewers can attempt to evade the detector or test benign edge cases and see the system respond in real time.

The interface is theme-aware (light/dark), responsive, and requires no external network access, so it runs fully offline on a laptop at a poster session.

5 Evaluation

We evaluate SHIELDMCP along three dimensions: attack-detection effectiveness, benign-task impact, and runtime latency. The benchmark harness comprises 40 MCP servers—25 benign servers drawn from popular open-source implementations (filesystem, database, web-search, code-execution, calendar) and 15 adversarial servers implementing the five threat families—with 487 attack scenarios and 200 benign task scenarios. The harness ships with the system and is reproducible with `python eval/run_evaluation.py --framework-only`.

Detection effectiveness. Table 1 reports attack success rate by family. Undefended, these attacks succeed at a weighted average of 70.9% across five LLM backends in the companion study. With the default heuristic backend, SHIELDMCP reduces overall ASR to **8.6%** (445/487 attacks detected); rug-pull and cross-tool-chain families are fully caught by the registry and correlation modules, while indirect prompt injection (16.4%) is the hardest for rules alone. Enabling the fine-tuned classifier backend detects all 487 attacks on this benchmark. Because the classifier training data and the evaluation scenarios are both synthetically generated and may share distributional structure, we treat the heuristic 8.6% as the conservative, deployment-realistic operating point and report the classifier backend as an optional higher-




Figure 2: The SHIELD MCP web interface running the *indirect prompt injection* scenario live. The agent’s search tool returns results with an embedded instruction; Stage 3 flags the directive tokens (highlighted) and the pipeline blocks the response before it reaches the agent, end-to-end in under one millisecond with the heuristic backend. The left panel lists the five attack families; the right panel shows the live benchmark and per-family attack success rate.

recall mode rather than the headline result.

Benign-task impact. On the 200 legitimate task scenarios, SHIELD MCP preserves a 95.0% task-completion rate (190/200 pass cleanly). The 5% drop comes from parameter sanitization conservatively flagging benign but unusual string patterns (e.g. a clipboard tool receiving shell-like text); this backend is rule-based and identical across the two detection modes. We regard this as an acceptable trade-off for security-sensitive deployments.

Latency. We measure per-call latency over 1,000 iterations on an Apple M5 Pro. The heuristic backend adds a median of **0.6 ms** end-to-end (P95 0.9 ms). With the DistilBERT classifiers loaded, median end-to-end latency is **17.2 ms** (P95 21.5 ms), dominated by the Stage-1 (8.9 ms) and Stage-3 (7.9 ms) neural passes. Since typical MCP tool calls themselves take hundreds of milliseconds, both backends add well under the 120 ms budget we target, and the heuristic backend is effectively free.

Expert review. To complement automated metrics, the companion study recruited three security engineers with production MCP experience to review 50 sampled alerts; evaluators agreed with the

ground-truth classification in 92% of cases, with disagreements concentrated on borderline descriptions containing legitimate instructional language. SHIELD MCP ships the review tool used for this study.

6 Comparison with Existing Systems

Guardrail systems such as LlamaFirewall (Chennabasappa et al., 2025) and NeMo Guardrails filter user inputs and model outputs but do not inspect the tool layer—the descriptions agents read, the parameters they construct, or the responses they ingest. Architectural approaches such as CaMeL (Debenedetti et al., 2025) offer strong control/data-flow guarantees but require substantial modifications to the agent pipeline that many production systems cannot accommodate. Training-time defenses (StruQ, SecAlign) harden a specific model but do not protect a heterogeneous fleet of servers. Recent MCP-specific tools—MCP-Guard (Xing et al., 2025), IPIGuard (An et al., 2025), and DRIFT—advance detection but are research prototypes rather than deployable, transport-level proxies. SHIELD MCP is distinguished by being simultaneously (i) modification-free and drop-in

at the MCP transport, (ii) covering all three interception points in one pass, and (iii) shipped as an installable package with an interactive interface and a reproducible benchmark, so its claims can be inspected rather than merely read.

7 Availability and Licensing

SHIELDMCP is released as open-source software under the Apache-2.0 license. The system is a Python package installable with `pip install shieldmcp`; the proxy (`shieldmcp proxy`), scanner (`shieldmcp scan`), and web demonstration (`shieldmcp demo`) are exposed as sub-commands. The repository includes the three-stage pipeline, the stdio and SSE proxies, the DistilBERT classifier training scripts and data generators, the full evaluation harness, and the web interface. The demonstration runs entirely locally with no API keys required; a hosted version is planned. Source, an installable package, and the screencast are linked from the project repository.

8 Conclusion

SHIELDMCP makes MCP tool-layer security concrete: a transparent, modification-free proxy that validates tool descriptions, parameters, and responses at runtime, packaged with an interactive interface that lets anyone launch an attack, watch it be blocked token-by-token, and probe the detector with their own inputs. By pairing a deployable system with a reproducible benchmark, we aim to help establish tool-layer security as a first-class concern as MCP adoption grows.

Ethical Considerations

This work studies and demonstrates security vulnerabilities in AI systems. All attack evaluations were conducted against locally deployed models and servers in controlled environments; no production systems or real users were targeted. The adversarial MCP servers we construct are for evaluation and are not novel zero-day disclosures—our threat taxonomy is derived from publicly available security research and the open-source SAFE-MCP project. We believe documenting these threats and providing defenses serves the public interest, since organizations cannot protect against risks they do not understand. The demonstration exposes only benign and clearly-labeled adversarial scenarios and does not facilitate the construction of new attacks beyond what is already public. We follow

the ACM Code of Ethics.

Limitations

SHIELDMCP has several limitations. The semantic and token classifiers are trained on English-language tool descriptions and synthetic data; multilingual and naturally-occurring adversarial servers would require additional training data, and the near-zero classifier ASR on our benchmark should be read with the caveat noted in Section 5 regarding synthetic-data overlap. The cross-call correlation module scales quadratically with the number of tool calls in a session, which can become a bottleneck in long-running interactions. Determined adversaries with knowledge of the defense can craft adaptive attacks; static detection is a component of, not a substitute for, defense-in-depth. Latency measurements were obtained on a single hardware configuration and will vary across deployments. Finally, our expert review involved three engineers, a sample that limits statistical power.

References

- Hengyu An, Jinghui Zhang, Tianyu Du, Chunyi Zhou, Qingming Li, Tao Lin, and Shouling Ji. 2025. Ipguard: A novel tool dependency graph-based defense against indirect prompt injection in llm agents.
- Sa-hana Chennabasappa, Cyrus Nikolaidis, Daniel Song, David Molnar, Stephanie Ding, Shengye Wan, Spencer Whitman, et al. 2025. Llamafirewall: An open source guardrail system for building secure ai agents. *arXiv preprint arXiv:2507.01814*.
- Edoardo Debenedetti, Ilia Shumailov, Tianqi Fan, Jamie Hayes, Nicholas Carlini, Daniel Fabian, Christoph Kern, Chongyang Shi, Andreas Terzis, and Florian Tramèr. 2025. Defeating prompt injections by design. *arXiv preprint arXiv:2503.12345*.
- Kai Greshake, Sahar Abdelnabi, Shailesh Mishra, Christoph Endres, Thorsten Holz, and Mario Fritz. 2023. Not what you’ve signed up for: Compromising real-world llm-integrated applications with indirect prompt injection. In *Proceedings of the 2023 Workshop on Artificial Intelligence and Security (AISec@CCS)*.
- Anthropic Salih et al. 2025. Model context protocol. *Linux Foundation*.
- Hao Song, Yiming Shen, Wenxuan Luo, Leixin Guo, Ting Chen, Jiashui Wang, Beibei Li, Xiaosong Zhang, and Jiachi Chen. 2025. Beyond the protocol: Unveiling attack vectors in the model context protocol ecosystem. *arXiv preprint arXiv:2507.01711*.

Zhiqiang Wang, Yichao Gao, Yanting Wang, Suyuan Liu, Haifeng Sun, Haoran Cheng, Guanquan Shi, Haohua Du, and Xiang-Yang Li. 2025. Mcptox: A benchmark for tool poisoning attack on real-world mcp servers. *arXiv preprint arXiv:2507.03234*.

Wenpeng Xing, Zhonghao Qi, Yupeng Qin, Yilin Li, Caini Chang, Jiahui Yu, Changting Lin, Zhenzhen Xie, and Meng Han. 2025. Mcp-guard: A defense framework for model context protocol integrity in large language model applications. *arXiv preprint arXiv:2507.04523*.

Weibo Zhao, Jiahao Liu, Bonan Ruan, Shaofei Li, and Zhenkai Liang. 2025. When mcp servers attack: Taxonomy, feasibility, and mitigation. *arXiv preprint arXiv:2507.07915*.